

Statistics refresher and linear regression

International Marketing Analytics



Dr Arne Floh, Senior Lecturer in Marketing
WU Vienna, Institute for International Business

arne.floh@wu.ac.at

[Office hours: book online \(Google Calendar\)](#)

Refresher

Python
by hand

OLS
theory

Tests

Chocolate
in Python

In class,
exercise

Refresher, Parts 1 to 5



- Scale types: what you may calculate
- Mean, median, standard deviation, standard error
- z-scores and the normal distribution
- Covariance and correlation
- All of it by hand in Python; quiz 1

Regression, Parts 6 to 11



- OLS, inference, fit, logs and elasticities
- Seven assumptions and four tests
- Regression in the real world
- statsmodels and plotnine on real data; quiz 2
- In-class assignment, coding exercise 1



Readings: Bojinov, Parzen & Hamilton (2025), Linear Regression (HBS); Skiera, Reiner & Albers (2022), Regression Analysis.

Our example: a chocolate kiosk

Five weeks, five prices

Week	Price (EUR)	Bars sold
1	2.00	110
2	2.50	100
3	3.00	80
4	3.50	70
5	4.00	40

One week, day by day

Day	Bars sold
Mon	10
Tue	12
Wed	11
Thu	13
Fri (event)	54



Small on purpose: every number today can be checked with a calculator, then in Python.

Scale types



Part 1 of 11 · nominal, ordinal, interval, ratio

Four scale types

Scale	What you can say	Kiosk example	Typical statistics
Nominal	same or different	brand, payment method	counts, mode
Ordinal	more or less	rating 1 to 5 stars, size S/M/L	median, percentiles
Interval	how much more (no true zero)	temperature in °C	mean, standard deviation
Ratio	how many times (true zero)	price, bars sold, revenue	all, incl. ratios



Two groups: nominal and ordinal are categorical; interval and ratio are metric (numbers you can average).

In regression: categorical variables enter as dummies (0/1), metric variables as they are.

Why the scale matters

20 °C is not twice as warm as 10 °C: no true zero

Mean of star ratings: treats ordinal as metric; say so

Store ID 17: a number, but only a label: nominal

Quick quiz: which scale?

- Country of the customer
- Satisfaction from 1 to 7
- Age in years
- Discount: yes or no
- Temperature on the day



Rule: check the scale before you average, chart or put a variable into a model.

Centre and spread



Part 2 of 11 · mean, median, variance,
standard deviation, standard error


Mean and median

mean $\bar{x} = \Sigma x / n = 100 / 5 = 20$
median middle of 10 11 12 13 54 = 12

Mean uses every value; one event day pulls it up

Median is robust: half the days below, half above

Use the median for skewed data and ordinal scales



Totals need the mean: 20 bars a day $\times$ 5 days = 100 bars in the week.

Variance and standard deviation, by hand

Week	Bars x	$x - \bar{x}$	$(x - \bar{x})^2$
1	110	30	900
2	100	20	400
3	80	0	0
4	70	-10	100
5	40	-40	1,600
Sum	400	0	3,000

Variance and standard deviation

$$\begin{aligned}\bar{x} &= 400 / 5 = 80 \\ s^2 &= 3,000 / (5 - 1) = 750 \\ s &= \sqrt{750} = 27.4 \text{ bars}\end{aligned}$$


Variance is in squared units (bars²)

Standard deviation is in bars: the typical distance from the mean



Why $n - 1$? deviations from the sample mean are a little too small; dividing by $n - 1$ corrects this. polars and pandas do it by default.

Same mean, different spread



3.8

bars: sd of kiosk A

31.6

bars: sd of kiosk B



For the manager: both sell 80 a week, but kiosk B needs much more safety stock and its forecasts are less certain.

Standard error: how precise is the mean?

$$SE = s / \sqrt{n}$$

$$95\% \text{ confidence interval} \approx \bar{x} \pm 1.96 \cdot SE$$

EUR 1.00

SE with 16 customers: $4 / \sqrt{16}$

EUR 0.50

SE with 64 customers: $4 / \sqrt{64}$

10.04-13.96

95 % interval for a mean
spend of EUR 12 ($n = 16$)

sd describes how much customers differ: spending varies by about EUR 4

SE describes how precise the average is: it shrinks with $\sqrt{n}$; four times the sample, half the SE



In every regression table: the column 'std err' is the standard error of a coefficient.

Standardisation and the normal distribution



Part 3 of 11 · z-scores, normal and standard normal distribution

Standardisation: z-scores

$z = (x - \bar{x}) / s$ a standardised variable has mean 0 and standard deviation 1


Kiosk	Bars this week	Mean	sd	z
Vienna	140	120	10	+2.0
Graz	75	60	15	+1.0
Our kiosk, week 5	40	80	27.4	-1.46



Read it: Vienna's week lies 2 standard deviations above normal, Graz's only 1: Vienna's week was the more unusual one.

Use: compare across stores, countries or units; spot outliers ($|z| > 3$); standardised coefficients.

The normal distribution



68 %

within ± 1 sd: 98 to 102 g

95 %

within ± 2 sd: 96 to 104 g


99.7 %

within ± 3 sd: 94 to 106 g



Why it matters: averages and regression coefficients are approximately normal (central limit theorem), so we can test them.

The standard normal distribution



z	share below z
1.00	84.1 %
1.645	95.0 %
1.96	97.5 %
2.576	99.5 %

Bars under 97 g? $z = (97 - 100) / 2 = -1.5$,
so 6.7 % of all bars



1.96: the number behind 95 % confidence intervals and two-sided tests at 5 %.

Covariance and correlation



Part 4 of 11 · do two variables move together,
and how strongly?

Covariance: do price and sales move together?



Week	$p - \bar{p}$	$x - \bar{x}$	product
1	-1.0	30	-30
2	-0.5	20	-10
3	0	0	0
4	0.5	-10	-5
5	1.0	-40	-40
Sum			-85

$$\text{cov} = -85 / (5 - 1) = -21.25$$




Sign = direction; the size depends on the units (EUR $\times$ bars).

Correlation: covariance without units

$$r = \text{cov} / (s_p \cdot s_x) = -21.25 / (0.79 \cdot 27.39) = -0.98$$

$$-1 \leq r \leq +1$$



Two warnings: r only measures straight-line relationships (the U has $r \approx 0$), and correlation is not causation.

Sources: Bojinov, Parzen & Hamilton 2025 (HBS note 9-622-100)

From correlation to regression

Not causation: ice-cream sales and sunburn rise together; sunshine drives both

Correlation says how strongly two variables move together, from -1 to $+1$

Regression says by how much y changes when x changes by one unit

slope $b = \text{cov}(p, x) / \text{var}(p) = -21.25 / 0.625 = -34$ bars per EUR
line through the means: $x = 80 - 34 \cdot (p - 3)$



Next: the least-squares line, built from exactly these pieces.

Python in Positron



Part 5 of 11 · mean, standard deviation, z-scores, covariance and correlation by hand, then checked

Step 1: the data and the mean

Python: polars

```
import polars as pl

kiosk = pl.DataFrame({
    "week": [1, 2, 3, 4, 5],
    "price": [2.0, 2.5, 3.0, 3.5, 4.0],
    "units": [110, 100, 80, 70, 40],
})

n = kiosk.height                # 5
mean_units = kiosk["units"].sum() / n  # by hand
kiosk["units"].mean()            # check
kiosk["units"].median()
```

Result	Value
n	5
mean, by hand	80.0
mean(), check	80.0
median()	80.0



Run it: click into a cell and press Ctrl/Cmd + Enter; the result appears in the Console.

Step 2: variance, standard deviation, standard error

Python: by hand, then checked

```
dev = kiosk["units"] - mean_units
dev          # 30, 20, 0, -10, -40

var_units = (dev ** 2).sum() / (n - 1)    # 750.0
sd_units = var_units ** 0.5              # 27.39
kiosk["units"].std()                     # check: n - 1

se_units = sd_units / n ** 0.5           # 12.25
```

Result	Value
variance	750.0
sd, by hand	27.39
std(), check	27.39
standard error	12.25



Note: std() divides by $n - 1$ (ddof=1), like the formula on the slide.

Step 3: standardise

Python: z-scores

```
kiosk = kiosk.with_columns(  
    z_units=(pl.col("units") - pl.col("units").mean())  
            / pl.col("units").std()  
)  
kiosk  
  
kiosk["z_units"].mean()    # 0  
kiosk["z_units"].std()     # 1
```

Week	Bars	z
1	110	1.10
2	100	0.73
3	80	0.00
4	70	-0.37
5	40	-1.46



Read it: week 5 lies 1.46 standard deviations below an average week.

Step 4: covariance and correlation

Python: by hand, then checked

```
dev_p = kiosk["price"] - kiosk["price"].mean()
dev_u = kiosk["units"] - kiosk["units"].mean()

cov = (dev_p * dev_u).sum() / (n - 1)          # -21.25
r = cov / (kiosk["price"].std()               # -0.981
        * kiosk["units"].std())

kiosk.select(
    pl.cov("price", "units").alias("cov"),
    pl.corr("price", "units").alias("r"))      # check
```

Result	Value
covariance	-21.25
correlation	-0.981
pl.cov, check	-21.25
pl.corr, check	-0.981




Same numbers as by hand: the computer only saves the typing.

Step 5: look at it

Python: plotnine

```
from plotnine import (ggplot, aes,  
                      geom_point, labs)  
  
(ggplot(kiosk, aes("price", "units"))  
 + geom_point(size=3)  
 + labs(x="Price (EUR)",  
        y="Bars sold per week"))
```



Plot first: a correlation can hide a curve, an outlier or two groups.

Your turn: by hand, then in Python

**Pairs, paper:
10 minutes**

**Python check:
10 minutes**

Day	Temp. °C	Bars
Mon	10	60
Tue	15	55
Wed	20	50
Thu	25	40
Fri	30	35

- 1 Scales:** which scale is each column?
- 2 Centre:** mean and median of bars sold
- 3 Spread:** variance and sd of bars, by hand
- 4 z:** z-score of the hottest day's sales
- 5 Together:** covariance and r of temperature and bars; then check all in Python

Quiz 1: statistics in five questions

1 Satisfaction from 1 to 5 stars is measured on which scale?

A nominal B ordinal C ratio

2 Daily sales 10, 12, 11, 13, 54: what describes a typical day best?

A the mean (20) B the median (12) C the standard deviation

3 You survey four times as many customers. The standard error of the mean ...

A halves B falls to a quarter C stays the same

4 A week's sales have $z = -2$. That week was ...

A 2 bars below the mean B 2 standard deviations below the mean C 2 % below the mean

5 Price and bars sold have $r = -0.98$. What does this mean?

A higher prices cause lower sales B opposite directions, nearly linear C 0.98 bars fewer per EUR



3 minutes alone, then compare with your neighbour. Same format as the online quiz and the exam.

Refresher: key takeaways

- 1 **Scale** first: it decides which statistics make sense
- 2 **Mean and median:** the median resists outliers; totals need the mean
- 3 **sd** describes the data; $SE = sd / \sqrt{n}$ describes the precision of the mean
- 4 $z = (x - \bar{x}) / s$ makes values comparable; 95 % of a normal lies within ± 1.96 sd
- 5 **Covariance** gives the direction, correlation the strength (-1 to $+1$), regression the size




Next: linear regression on real chocolate scanner data, Parts 6 to 11.

Linear regression: theory



Part 6 of 11 · real data, simple and multiple regression, inference, fit, logs, from elasticity to price

The question: what does a price cut sell?



Weekly sales of brand 1; blue: weeks with a price cut (price below 1.20).

Data: 68 weeks of real scanner data, 4 brands

Normal week: price about 1.50, sales about 100

Price-cut week: price about 1.03, sales up to 530

Also in the data: competitor prices, features, displays, temperature, holidays



The manager asks: how much does 1 % lower price sell, and is the regular price right?

Correlation on real data: how strong is strong?

Rule of thumb for $|r|$

$ r $	Interpretation
0 to 0.2	very weak
0.2 to 0.4	weak to moderate
0.4 to 0.6	medium to substantial
0.6 to 0.8	very strong
0.8 to 1.0	extremely strong

−0.94 sales and own price of brand 1:
an extremely strong negative
relationship

Kiosk: $r = -0.98$, five made-up weeks (Part 4)


Both: extremely strong, yet far from a perfect
line



Real data are never a perfect line: competitors, features and the season move sales too.

Sources: Bojinov, Parzen & Hamilton 2025 (HBS note 9-622-100)

Simple linear regression



$$\text{sales} = b_0 + b_1 \cdot \text{price1} + e$$

$b_0 = 1,110$: intercept, sales at price 0
(not meaningful here)


$b_1 = -669$: 10 cents more price, about 67
units fewer a week

e : residual, what the line does not explain



Look at the chart: the points curve, so a straight line misses the price-cut weeks.

Least squares: the line with the smallest squared errors



residual: $e_i = y_i - \hat{y}_i$
OLS: minimise $\sum e_i^2$

Squared: misses above and below do not cancel

Large misses count much more than small ones

One line: the unique best fit for these data

Sources: Skiera, Reiner & Albers 2022; Bojinov, Parzen & Hamilton 2025

Multiple regression: holding other things constant

$$\ln \text{ sales} = b_0 + b_1 \cdot \ln \text{ price1} + b_2 \cdot \ln \text{ price2} + b_3 \cdot \ln \text{ price3} + b_4 \cdot \ln \text{ price4} \\ + b_5 \cdot \text{feature1} + b_6 \cdot \text{temp} + b_7 \cdot \text{december} + e$$



Own mix

Price of brand 1
Feature ads, displays



Competitors

Prices of brands 2 to 4
Their features and displays



Context

Temperature
December and Easter weeks



Each b: the effect of one variable with all others held constant (*ceteris paribus*).

Left out a relevant variable? Its effect ends up in the others (omitted variable bias).

Is the effect real? t-test, p-value, interval, F-test

t-test

$$t = b / se$$

Own price: $-3.64 / 0.30 = -12.1$

p-value

Below 0.05: significant

Own price: $p < 0.001$; brand 4: $p = 0.86$

95 % confidence interval

About $b \pm 2 se$

Own price: -4.24 to -3.04

F-test

Are all slopes together zero?

$F = 59.5$, $p < 0.001$



Significant is not important: read the size of b and its interval, not just the stars.

Sources: Bojinov, Parzen & Hamilton 2025; Skiera, Reiner & Albers 2022

How well does the model fit?

0.874

R^2 : the model explains 87 % of the weekly variation in log sales

0.859

adjusted R^2 : R^2 with a penalty for each extra variable

0.169

RMSE: the typical weekly miss, about 17 % in a log model

Careful



- R^2 never falls when you add a variable
- Compare R^2 only for the same outcome (sales vs log sales)
- High R^2 does not mean causal



Main model for brand 1: log-log, own and competitor prices, feature, temperature, December; 68 weeks.

Sources: Bojinov, Parzen & Hamilton 2025; Skiera, Reiner & Albers 2022

Functional forms and dummy variables

Model	Equation	b means	Brand 1
Linear	$\text{sales} = a + b \cdot \text{price}$	+1 unit price: b units	$b = -669$ units
Log-log	$\ln \text{sales} = a + b \cdot \ln \text{price}$	+1 % price: b %	$b = -3.64$ (elasticity)
Log-linear	$\ln \text{sales} = a + b \cdot \text{temp}$	+1 °C: about $100 \cdot b$ %	$b = -0.0066$: -0.66 %
Dummy in log model	$\ln \text{sales} = a + b \cdot \text{december}$	$e^b - 1$, vs other weeks	$b = -0.385$: -32 %



Why logs? sales respond in per cent, not in units, and the elasticity can be compared across brands and countries.

Sources: Skiera, Reiner & Albers 2022; Wooldridge 2025, ch. 6 to 7

From elasticity to a pricing decision

optimal price $p^* = \text{marginal cost} \times \varepsilon / (1 + \varepsilon)$ (valid for $\varepsilon < -1$)

\$52

Skiera et al.: $\varepsilon = -2.34$, unit cost \$30 $\rightarrow p^* = 30 \times 2.34 / 1.34$

$\times 1.38$

Brand 1: $\varepsilon = -3.64 \rightarrow$ price 38 % above marginal cost

Short term: promotion weeks inflate the elasticity; regular-price elasticities are smaller

Competitors react, and the retailer takes its margin

Sources: Skiera, Reiner & Albers 2022; Hanssens, Parsons & Schultz 2001

Assumptions and tests





Part 7 of 11 · what OLS assumes, how to look,
how to test, what to do

Seven assumptions behind OLS

Assumption	If violated	Look or test (statsmodels)
1 Linear form	biased coefficients	residuals vs fitted; try logs
2 Exogeneity, no omitted variables	biased, not causal	theory, controls, experiments
3 Constant error variance	wrong standard errors	funnel in plot; Breusch-Pagan
4 No autocorrelation	standard errors too small	residuals over time; Durbin-Watson ≈ 2
5 Normal residuals	tests unreliable in small samples	histogram; Jarque-Bera
6 No multicollinearity	unstable b, large se	correlations; VIF below 10 (better 5)
7 Enough data, no dominant week	fragile results	$n \geq 3 \text{ to } 5 \times \text{coefficients}$; Cook's distance

Sources: Skiera, Reiner & Albers 2022; Hair et al. 2019; Wooldridge 2025

Look first: two residual plots



Look for: a curve (wrong form) or a funnel (unequal variance)

Look for: a rough bell shape, no extreme weeks



Brand 1: no curve, no funnel, a roughly bell-shaped histogram.


Then test: four numbers for the brand 1 model

1.41 Durbin-Watson: below 2, some autocorrelation between weeks

0.14 Breusch-Pagan p: no evidence of unequal variance

0.07 Jarque-Bera p: residuals roughly normal

2.4 largest VIF: no multicollinearity



Verdict: usable. Report Newey-West (HAC) standard errors because of the autocorrelation: 0.32 instead of 0.30 for own price.

When a test fails: what to do

Problem	Remedy	In statsmodels
Curve in residuals	logs, squares, interactions	<code>np.log(x), l(x**2)</code>
Unequal variance	robust standard errors	<code>fit(cov_type="HC3")</code>
Autocorrelation	HAC standard errors; lags; ARIMA errors (Session 4)	<code>fit(cov_type="HAC", cov_kwds={"maxlags": 1})</code>
Multicollinearity	more variation, combine variables, experiments	<code>variance_inflation_factor</code>
Dominant weeks	check data, report with and without	<code>get_influence().cooks_distance</code>
Endogeneity	experiments, instruments, fixed effects	Sessions 2 to 5

Sources: Skiera, Reiner & Albers 2022; statsmodels documentation

Three warnings from real data

Displays



- Only 4 display weeks
- VIF 11.4 and 10.9
- Effects cannot be separated

One Easter week



- A dummy for one week fits it exactly
- Leverage 1.00
- No real estimate of Easter

Mailings



- Skiera et al.: mailings and salespersons $r = 0.989$
- VIF 53
- Salesperson effect flips sign



The real fix: vary the marketing instruments independently, ideally in a field experiment.

Sources: Skiera, Reiner & Albers 2022; own calculation on the chocolate data

Endogeneity: when marketing follows demand

The assumption: explanatory variables are unrelated to the error term

Marketing breaks it: budgets go where managers expect success, so spend and unexplained sales rise together

Brand 1: price cuts come with features, so the price effect partly measures the feature

Evidence: in 15 Facebook experiments, regression often overstated the true lift, sometimes 3 times



Remedies: experiments (A/B, geo tests), instrumental variables such as cost shifters, panel data with fixed effects.

Sources: Skiera, Reiner & Albers 2022; Gordon, Zettelmeyer, Bhargava & Chapsky 2019

Regression in the real world



Part 8 of 11 · benchmarks, companies, what transfers to MMM

Benchmarks from decades of regressions

-2.62 mean price elasticity, 1,851 estimates (Bijmolt et al. 2005)

0.12 mean short-term advertising elasticity (Sethuraman et al. 2011)

-0.42 mean price elasticity in 14 Asia-Pacific countries, lower with high power distance (Datta et al. 2022)

-3.64 brand 1, chocolate: promotion-driven, many substitutes



Use benchmarks to judge whether your own estimate is plausible.

Where companies use regression

Meta Robyn

MMM in production:
ridge regression on media

Booking.com

With Microsoft: CUPED,
regression makes A/B tests
faster

Google

Geo experiments: test vs
control regions by regression

Walmart (M5)

Sales forecasts need price,
promotion and event
regressors

statworx

Frankfurt consultancy: price
elasticities, log-log

Bayer

Budget across countries
from estimated elasticities



Session 2: the marketing mix model is this regression with adstock and saturation.

Sources: Robyn docs; Deng et al. 2013; Vaver & Koehler 2011; Makridakis et al. 2022; Fischer et al. 2011

The chocolate data in Python



Part 9 of 11 · load and plot, the first regression,
the log-log model, the checks

The data: 68 weeks of chocolate scanner data

Column	Meaning
week	1 to 68
sales	weekly unit sales of brand 1
price1 to price4	prices of brands 1 to 4
feature1 to feature4	share of stores with a feature ad, 0 to 1
display1 to display4	share of stores with an in-store display
fand1, fand3, fand4	share with feature and display together
temp	temperature, °C
december, easter	1 in December and Easter weeks

File: data/legacy/chocolate_dataset.csv


Step 1: load and look

Python: polars and plotnine

```
import polars as pl
from plotnine import (ggplot, aes, geom_point,
                      geom_smooth, labs)

choc = pl.read_csv(
    "data/legacy/chocolate_dataset.csv")
choc.select("sales", "price1",
            "price2").describe()

(ggplot(choc, aes("price1", "sales"))
 + geom_point()
 + geom_smooth(method="lm")
 + labs(x="Price of brand 1",
        y="Weekly sales"))
```



Look before you model: ranges, outliers, the shape of the relationship.

Step 2: the first regression

Python: statsmodels

```
import statsmodels.formula.api as smf

df = choc.to_pandas()
simple = smf.ols("sales ~ price1",
               data=df).fit()
print(simple.summary())

simple.params      # b0, b1
simple.rsquared    # 0.884
```

	coef	std err	P> t
Intercept	1,110.0	44.0	0.000
price1	-668.8	29.8	0.000

$R^2 = 0.884$, $n = 68$

Read it: 10 cents more price, about 67 units fewer per week; the price explains 88 % of the weekly variation.



But: the residuals curve. Next: the log-log model with the marketing mix.

Step 3: the log-log model with the marketing mix

Python: statsmodels

```
import numpy as np

model = smf.ols(
    "np.log(sales) ~ np.log(price1)"
    " + np.log(price2) + np.log(price3)"
    " + np.log(price4) + feature1"
    " + temp + december",
    data=df).fit()
print(model.summary())
```

	coef	p
ln price1	−3.64	0.000
ln price2	−0.80	0.076
ln price3	0.36	0.109
ln price4	−0.03	0.855
feature1	0.18	0.307
temp	−0.007	0.045
december	−0.38	0.009

$R^2 = 0.874$, adjusted 0.859, n = 68



Read it: 1 % higher price, 3.6 % fewer sales; December weeks 32 % lower; each degree warmer 0.7 % lower.

Step 4: check the assumptions

Python: tests and robust errors


```
from statsmodels.stats.stattools import (
    durbin_watson, jarque_bera)
from statsmodels.stats.diagnostic import het_breuschpagan
from statsmodels.stats.outliers_influence import (
    variance_inflation_factor)

durbin_watson(model.resid)                # 1.41
het_breuschpagan(model.resid, model.model.exog)[1] # 0.14
jarque_bera(model.resid)[1]                # 0.07
X = model.model.exog
[variance_inflation_factor(X, i) for i in range(1, 8)]

hac = model.model.fit(cov_type="HAC",
                      cov_kwds={"maxlags": 1})
```

Residual plot

```
from plotnine import geom_hline
df["fit"] = model.fittedvalues
df["res"] = model.resid
(ggplot(df, aes("fit", "res"))
 + geom_point()
 + geom_hline(yintercept=0))
```



Decide: report the HAC result; the own-price elasticity stays -3.64 , its standard error rises to 0.32 .

Quiz 2: regression in five questions

1 What does OLS minimise?

A the sum of residuals B the sum of absolute residuals C the sum of squared residuals

2 In $\ln \text{sales} = a + b \cdot \ln \text{price}$, $b = -3.64$ means ...

A 1 EUR more, 3.64 units fewer B 1 % higher price, 3.6 % fewer sales C 3.64 fewer per week

3 A VIF of 11 signals ...

A autocorrelation B multicollinearity C heteroscedasticity

4 Weekly data, Durbin-Watson = 1.41. This suggests ...

A positive autocorrelation B no problem C unequal variance

5 R^2 rises from 0.85 to 0.86 after adding a variable. This shows ...

A the variable matters B nothing: R^2 never falls C the model is causal



3 minutes alone, then compare with your neighbour. Same format as the online quiz and the exam.

In-class assignment



Part 10 of 11 · groups of 3, 25 minutes, on the chocolate data

In-class assignment: pricing brand 1

**Groups of 3:
20 minutes**

**Plenary:
5 minutes**

- 1 Competition:** plot sales against price2, price3 and price4. Which brand competes most?
- 2 Displays:** add display1 and fand1 to the log-log model. What happens to the VIFs and coefficients, and why?
- 3 Price:** marginal cost is 0.80. What price does $\epsilon = -3.64$ suggest? Why is the real price about 1.50?
- 4 The manager:** write one sentence for the brand manager, with a number and one caveat.

File: sessions/01-foundations/regression_chocolate.qmd (the code of Part 4 and these questions)

Coding exercise 1



Part 11 of 11 · individual, 5 % of the grade

Coding exercise 1: what to do

- 1 Explore:** a `great_tables` summary and two plotnine charts: sales over time, sales against price1
- 2 Simple model:** $\text{sales} \sim \text{price1}$; interpret the slope for a 10-cent price change
- 3 Log-log model:** own and competitor prices, feature1, display1, temp, december; interpret 3 coefficients
- 4 Assumptions:** residual plot, histogram, Durbin-Watson, Breusch-Pagan, Jarque-Bera, VIF; one sentence each
- 5 Robust errors:** refit with HAC standard errors; what changes and what does not?
- 6 Recommend:** the optimal price for marginal cost 0.90, with two caveats, in at most 100 words

Brief: `exercises/exercise-1.qmd` in the course repository and on Canvas

Coding exercise 1: submission and grading

Submit on Canvas



- exercise1_lastname.qmd and the rendered .html
- Renders from top to bottom without errors
- Every number in the text from code
- Prompt log at the end: what you asked, what you changed
- Deadline: see Canvas

Criterion	Points
Code runs, charts and tables clean	5
Interpretation of the coefficients	6
Assumption checks and conclusion	5
Recommendation with caveats	2
Prompt log	2
Total (5 % of the grade)	20



Individual work: discuss ideas, but write and check your own code. AI assistants are allowed; you own every number.

Regression: key takeaways

- 1 **OLS** draws the line with the smallest squared errors; each b holds the other variables constant
- 2 **Logs** turn coefficients into elasticities: brand 1, -3.64
- 3 **Look, then test:** residual plots, Durbin-Watson, Breusch-Pagan, Jarque-Bera, VIF
- 4 **Fix what fails:** robust or HAC errors, better form, more variation; experiments for causality
- 5 **Decide:** $p^* = \text{cost} \times \varepsilon / (1 + \varepsilon)$, with the caveats



Next session: the marketing mix model, a regression with adstock and saturation.

References: statistics refresher

Books and readings



Field, A. (2024). Discovering Statistics Using IBM SPSS Statistics, 6th ed. SAGE. Chapters 1 to 2 and 8.

Hair, J. F. et al. (2019). Multivariate Data Analysis, 8th ed. Cengage. Chapter 2.

Stock, J. H. & Watson, M. W. (2019). Introduction to Econometrics, 4th ed. Pearson. Chapters 2 to 3.

Online



Seeing Theory. Brown University: interactive visual introduction to probability and statistics (free).

StatQuest. Videos on the normal distribution, standard deviation vs standard error, covariance and correlation (free).

polars documentation. mean, median, std, cov, corr.

Gelman, Hill & Vehtari (2020). Regression and Other Stories, ch. 3 to 4 (free PDF).

References: linear regression

Academic



Bojinov, I. I., Parzen, M. & Hamilton, P. J. (2025). Linear Regression. HBS note 9-622-100 (Canvas).

Skiera, B., Reiner, J. & Albers, S. (2022). Regression analysis. In Homburg et al. (eds), Handbook of Market Research, 299–327. Springer.

Wooldridge, J. M. (2025). Introductory Econometrics, 8th ed. Cengage.

Hanssens, D. M., Parsons, L. J. & Schultz, R. L. (2001). Market Response Models. Kluwer.

Bijmolt, T., van Heerde, H. & Pieters, R. (2005). Journal of Marketing Research 42(2).

Gordon, B. R. et al. (2019). Marketing Science 38(2).

Learning resources



James, G. et al. (2023). ISLP, ch. 3 and the Python lab (free).

Schwarz, Chapman & Feit (2020). Python for Marketing Research and Analytics, ch. 7. Springer.

statsmodels. Documentation and example notebooks.

Turrell, A. Coding for Economists: regression and diagnostics (free).

StatQuest. Linear regression, clearly explained (video).

DataCamp. Introduction to Regression with statsmodels in Python.



VIENNA UNIVERSITY OF
ECONOMICS AND BUSINESS

**INSTITUTE FOR INTERNATIONAL
BUSINESS**

Welthandelsplatz 1, 1020 Vienna,
Austria

DR ARNE FLOH

Senior Lecturer in Marketing

F +43-31336-6367

arne.floh@wu.ac.at

www.wu.ac.at/en/welthandel/arne-floh

[Office hours: book online](#)